

MAËLYS • HEAD OF AI • HOME ASSIGNMENT

Creative *Loop*


An AI-powered, self-improving acquisition engine — a closed feedback loop between creative production and media performance.

Part 1 • System design

Part 2 • Working prototype: video ad → landing page → next creative

WHERE THE MONEY LEAKS TODAY

Nine steps, each with AI inside — and nothing learning between them



↑ every hand-off is manual · every learning stays in one person's head or one tool's archive

Velocity is bounded by hours

Manual hand-offs cap testing volume — the raw material of learning.

Learnings evaporate

Why an ad won lives in one buyer's head. New briefs start from taste.

Message match dies at the click

Your Caroline LP = the generic page except H1 + hero. Her promise is gone one scroll in.

The compounding asset isn't the ad. It's the learning.


Everyone can generate infinite creative. The brands that win **learn fastest per dollar of spend**. The engine's job: **make every dollar teach the system what to make next.**



Five layers, one loop

CREATIVE INTELLIGENCE — the Creative Genome

Every asset (ad, LP, quiz) is auto-tagged at birth with a structured taxonomy: hook type · angle · persona · format · creator · claims · offer framing · LP archetype. Multimodal models watch every video; tagging is applied retroactively to the entire historical library.



GOVERNANCE — humans stay where humans win

Compliance gate auto-checks claim hedging ("the look of") before anything ships · creative-director approval on new angles · media buyers own overrides. AI drafts; people direct.

The Creative Genome: measure attributes, not ads

One ad's ROAS is an anecdote. Spend measured across a shared taxonomy is evidence — starting with your **historical library**, before a single new ad is made.

GENE	EXAMPLE VALUES (FROM YOUR REAL CREATIVE)	WHY IT PAYS
Hook type	creator confession · problem callout ("GLP-1 loose skin") · result tease · question	Hooks drive thumbstop — the cheapest KPI to move
Angle	post-weight-loss skin · works-while-you-sleep · skeptic conversion	Angles fatigue as a group; detect it before ROAS shows it
Format · Creator	UGC talking head · explainer · montage · Caroline B. · AI-generated	Creator × angle fit is your scaling constraint
Offer framing	TBYB "\$0" · "\$5 trial" · "or don't pay"	Framing moves CVR <i>and</i> refund rate — must be tracked to LTV
LP archetype	advertorial · compact offer-first · creator listicle · quiz	Ad→LP pairing is a free CVR lever nobody measures

Retroactive tagging = instant history

Models watch every historical ad once; tags join spend you already paid for.

Attribution stays honest

Learnings are priors. The system proposes; the test confirms.

What flows in, what flows out

INPUTS

SOURCE	FEEDS
Historical creatives + spend	The engine's starting memory
Meta insights (CPM, thumbstop, CTR, CVR, CPA, AOV, ROAS)	Attribute learnings; scale/retire rules
GA4 site behaviour	Section engagement → page iteration
LTV + 1st-party data	LTV-adjusted targets; refund guardrails
A/B test archive	Seeded creative memory
Ads Library + competitor sites	Angle-gap map

OUTPUTS

ARTIFACT	CONSUMED BY
Learning-grounded briefs & scripts	Creative team & creators
AI video teasers (Veo)	Cheap kill-fast hook pretests
Hook-matched landing pages	Web/CRO — Part 2 does this today
Weekly Creative Intelligence report	Growth team + leadership
Launch/retire recommendations	Media buyers (with override)
Creative memory (hypothesis + result)	The next cycle — the loop itself

Four concrete feedback mechanisms

1 · LEARNINGS IN THE PROMPT

The brief generator is prompted with the live learning table — *"creator-confession hooks +35% ROAS."* Every concept must cite a learning, or declare itself an exploration.

2 · EXPLORE / EXPLOIT BUDGET

70% iterate winners · 20% new hooks on proven angles · 10% swings. Enforced by the concept queue itself.

3 · FATIGUE TRIGGERS REFRESH

CTR decay per **angle**, not per ad, auto-triggers refresh briefs — before decay reaches ROAS.

4 · OPTIMIZE TO LTV, NOT FIRST PURCHASE

Winners re-ranked on 90-day LTV:CAC and refund rate — a "\$0" frame can win CVR while leaking into returns.

Learn from history first. Generate second. Automate last.

Sequenced by risk: each phase pays for itself before the next adds autonomy.

CRAWL

Retroactive intelligence

- Genome auto-tagger on the historical library
- Join tags ↔ Meta · GA4 · LTV · A/B archive
- First weekly intelligence report

Zero workflow change. Insight from money already spent.

WALK

Learning-grounded generation

- Brief & script generator, learnings in context
- Ad→LP builder in production (Part 2)
- Veo teaser pretests
- Human approval gates on everything

Velocity ×3–5 at constant headcount.

RUN

Closed-loop automation

- Launch/test/retire rules via Meta API
- Bandit-style explore/exploit allocation
- LTV-objective optimization · LP personalization

The loop runs nightly.

THE FIRST FIVE BUILDS

Ordered by evidence-per-dollar, not by glamour

#	BUILD	WHY THIS, WHY NOW	PAYOFF
1	Creative Genome + retro-tagger	Turns your ad history + A/B archive into structured evidence. Everything else consumes this.	The engine's memory
2	Learning engine + weekly report	Makes the intelligence visible and challengeable from day one.	Trust + direction
3	Brief & script generator	Concepts grounded in learnings, in MAËLYS voice, hedged by construction. Creators create; the engine aims.	Creative velocity
4	Ad→LP builder PROTOTYPE BUILT	Message match is the biggest post-click CVR lever — this is Part 2, working today.	CVR on every ad
5	Iteration & pretest engine	Winner variations + Veo teasers kill weak hooks for dollars, not production budgets.	Cheaper, faster testing

Measure the engine, the business, and the failure modes

ENGINE HEALTH

- **Creative hit rate** — % beating portfolio-median ROAS at the test gate
- **Concept→live time** — queue to first impression, in days
- **Learning velocity** — validated learnings per month
- **Cost per winning creative** — test spend ÷ winners

BUSINESS

- **Blended CAC & ROAS** — Meta cross-checked vs GA4 + first-party
- **90-day LTV:CAC by creative cohort** — the scaling decision
- **New-creative share of spend** — aging-winner risk
- **ATC & LP CVR per archetype** — GA4, per generated page

GUARDRAILS

- **Refund rate per offer framing** — "\$0" must not buy CVR with returns
- **Fatigue index** — CTR decay slope per angle
- **Compliance flag rate** — 0 unhedged claims shipped
- **Brand consistency** — weekly sampled human QA

I built the loop — it runs today, and it shows its sources

Three stages, all live: ad in → analysis → landing page → a new AI-generated ad. Every listed data source consumed — and labeled where it's used.

1 · Ad Intelligence

Gemini watches the full video and returns its creative DNA — hook, angle, persona, claims, **compliance flags**, recommended archetype.

2 · Hook-matched landing page

Your real archetypes, voice, offer and **live CDN imagery** (likeness-gated) + optional free-text direction. Every section carries a **provenance chip**.


Data as fuel, not scoreboard

Fixed internal dataset + **4 competitor funnels scraped live** — cited sources with clickable links. The new page's own results are unknown until it runs.

3 · The loop closes — with Veo

Learnings + competitor gaps → three concepts (safe / new-hook / swing) → one click renders a real 8-second Veo ad.

Grounded generation, not blind generation



The test queue I'd run first

PRI	TEST	HYPOTHESIS	METRIC
1	Deep hook-match on creator pages — full creator variant, auto-generated (the Part 2 builder does this)	Creator voice past the hero lifts ATC for creator traffic	ATC / CVR
2	Performance pass — async VWO, lazy-load below fold, poster-frame videos	Faster LCP recovers the paid clicks currently lost pre-first-paint	LP CVR, CPC efficiency
3	One price frame — "\$5 shipping · 3 full weeks" everywhere; pre-select an offer option	Coherence beats the "\$0" trick: same CVR, less refund drag	CVR + refund rate
4	Sticky ATC on the long advertorials (module exists on the listicle)	19-section pages convert mid-scroll readers they currently strand	ATC
5	Quiz keeps its promises — echo the user's answers in the result + restate the \$35 gift	Perceived personalization converts the yes-ladder's commitment	quiz→LP CTR, CVR
6	Consistency sweep — one jars-sold number, one clinical footnote, real scarcity or none	Cross-LP shoppers stop catching contradictions that undercut the clinical positioning	trust · brand QA
7	Post-trial "empty-jar" guarantee — 60-day money-back after the TBYB charge (The Body Firm's signature; no MAËLYS LP has it)	Extending risk reversal past the trial charge closes the objection TBYB leaves open	CVR + refund rate

CLOSE

In 90 days: the engine is
learning from every dollar
you've ever spent — and
writing the next brief.

Crawl · Genome + retro-tagging + first intelligence report

Walk · learning-grounded briefs + LP builder in production

Run · Veo pretests + automated test cells

The product works while she sleeps. The engine should too.